



arm

Arm server training armv8.x architecture

Zenon Xiu
Senior Applications Engineer Manager
Partner Enablement Group

Taipei, July 2018

Copyright © 2018 Arm Limited

About me

Zenon Xiu (修志龍)

- Joined arm Partner Enablement Group at 2009
- Based in Shanghai
- Support arm partners on



TrustZone®
System Security by ARM



arm

What PEG does --Arm Services

Support Services

Training

- Expand employee's knowledge and capabilities
- Reduce time to market
- Fully exploit product features

Technical Support

- Worldwide support
- Design efficiently and effectively
- Errata updates and delivery

Documentation

- Accessible product information
- Comprehensive and up-to-date
- Supplementary online resources

Arm Design Reviews

- End-to-end project assistance
- Optimize subsystem designs
- Reduce risk


www.arm.com/support

ARM Training Overview

ARM Training gets partner teams up to speed quickly

- Detailed training across all ARM Products
- Covers advanced technologies such as ARM® DynamIQ and ARM® TrustZone®

Delivered by senior Arm engineers in a way that works best for you and your team




Face to face private training
delivered by experienced ARM
trainers at a location of your choice
Private courses can be customizable
to meet your specific needs



Live virtual classroom delivered by
an experienced ARM trainer
Private and public courses available
Delivered to you, wherever you are
at a scheduled date and time




Online learning - short topics
accessible wherever you are and
whenever you need them.
To be launched Oct 2017



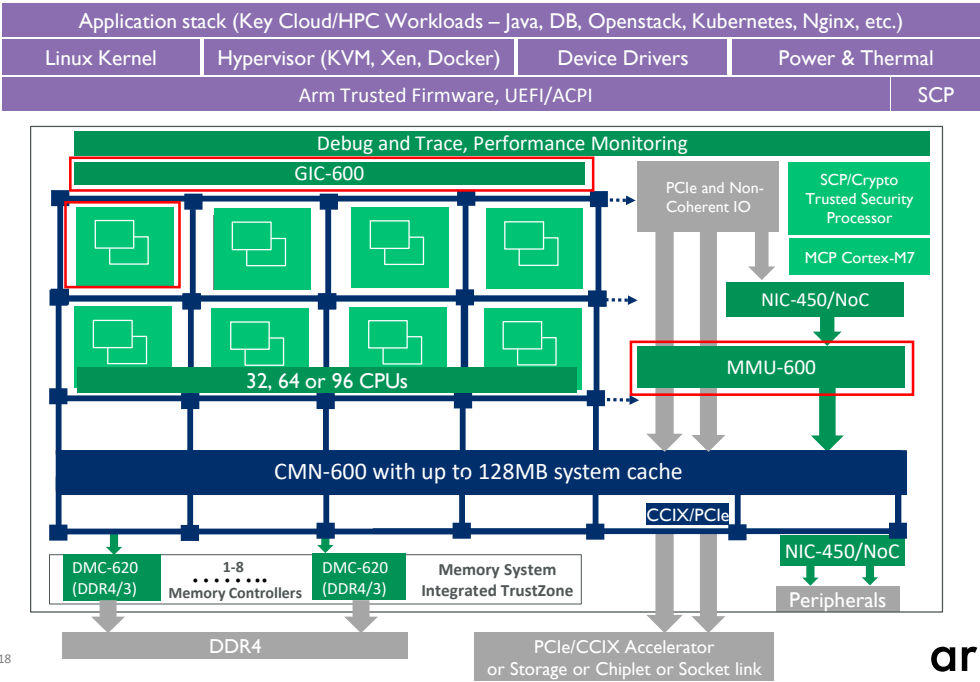
Arm Services

Providing the help you need, when you need it.

A range of **support services** are available to help reduce risk and reduce your time to market.



www.arm.com/support



Agenda

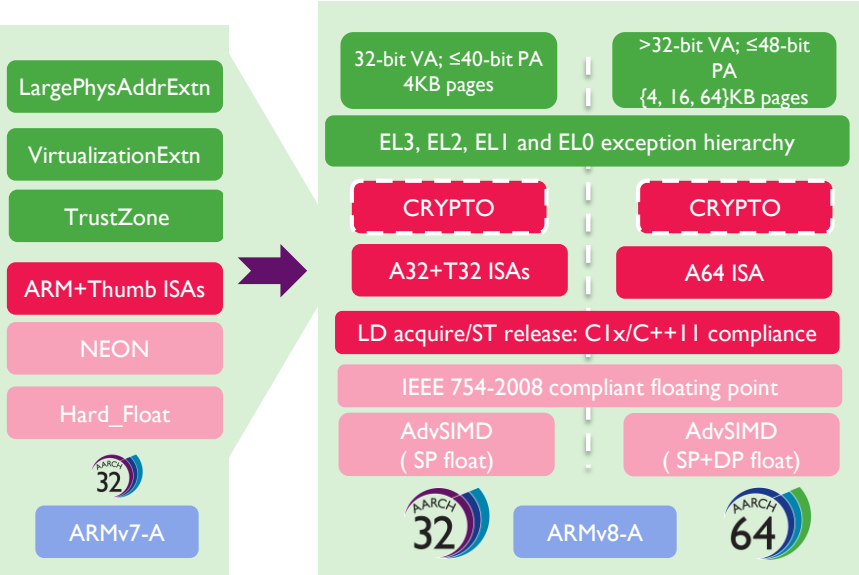
Armv8.x overview

- Registers
- ISA
- SVE
- Exception
- Memory management
- Memory model
- Atomic
- Architecture timer
- Virtualization
- MPAM
- Security
- RAS

7 Confidential © Arm 2018



Arm Architecture Evolution: Armv8.0-A



8 Confidential © Arm 2018

8

arm

AArch64

New 64-bit general purpose registers (X0 to X30)

New instructions – A64, fixed length 32-bit instruction set

- Includes SIMD, floating point and crypto instructions

New exception model

- EL0-EL3

Virtual addresses now stored in 64-bit registers

- Extend virtual memory space

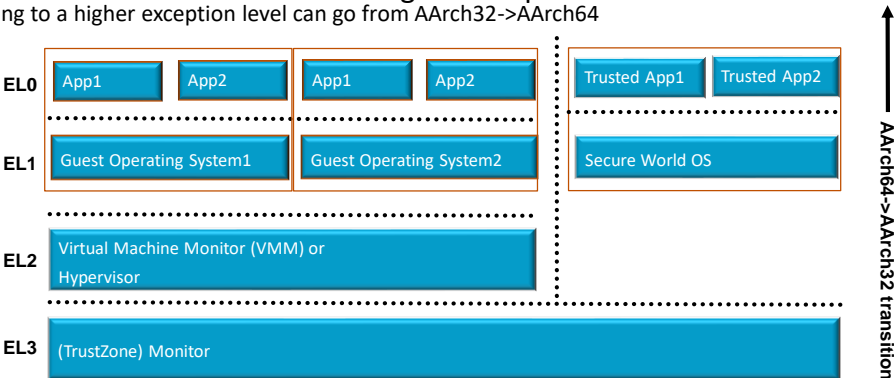
Armv8 Exception Model

4 exception levels: EL3-EL0

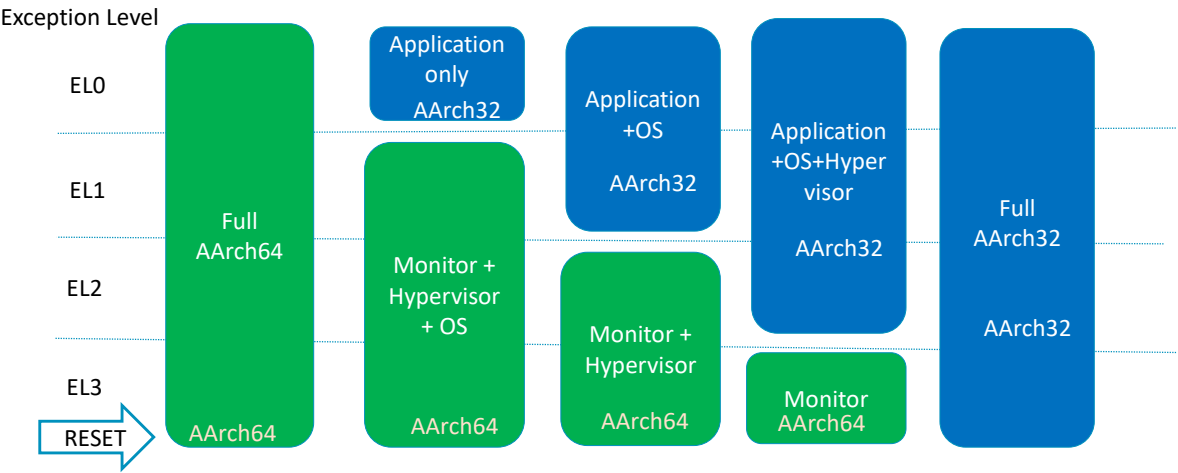
- Forms a privilege hierarchy, EL0 the least privileged

Exceptions can be taken to same or a higher exception level

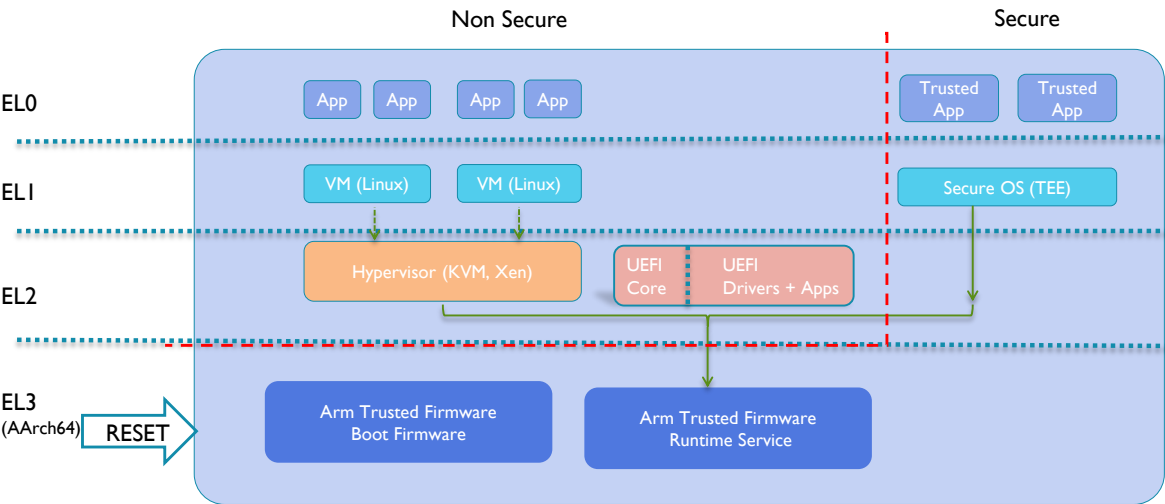
- Going to a higher exception level can go from AArch32->AArch64



ARMv8 AArch64/AArch32 transaction



Arm v8 software stack example



Armv8.x Architecture Features

V8.0-A

64-bit architecture support
64-bit HW support for Virtualization
SIMD + FP
Cryptography support for SHA1, SHA256, AES.
CRC instructions

Memory model
virtualization
FP & SIMD
Crypto & security
Monitoring & debug

v8.1-A

VHE: Improved support for Type-2 Hypervisors (KVM)
Rounding Double Multiply Add/Subtract
Atomic Instructions
Limited Order Regions
Hardware update of Access & Dirty Bits

v8.2-A

Half-Precision support
RAS
Statistical Profiling

v8.3-A

Improved support for nested Hypervisors
JavaScript conversion function
Complex Number support
Pointer Authentication
RCpc load/store

v8.4-A

Further support for nested virtualization.
Secure EL2.
16x16+32 FP.
Dot Product.
SHA512, SHA3
SM3, SM4
Timing Insensitivity of DP Instructions
Changes to memory model
Small page table support
Activity Monitors
MPAM

13 Confidential © Arm 2018

arm

Agenda

Armv8.x overview

Registers

ISA

SVE

Exception

Memory management

Memory model

Atomic

Architecture timer

Virtualization

MPAM

Security

RAS


14 Confidential © Arm 2018



AArch64 Registers

AArch64 provides 31 general purpose registers: R0-R30


- Each register has a 64-bit (**Xn**) and 32-bit (**Wn**) form, common in all ELs



AArch64 introduces the “zero” register: XZR and WZR

Separate register file for floating point, SIMD, and crypto operations: V0-V31

- Each register has a 128-bit (**Qn**), 64-bit (**Dn**), 32-bit (**Sn**), 16-bit (**Hn**), and 8-bit (**Bn**) form



15 Confidential © Arm 2018

arm

AArch64 Registers

AArch64 provides banked SP registers for each EL

AArch64 has PSTATE bit fields (not a register) to indicate or control program status


- Bit field could be changed individually, e.g.



- Access PSTATE
 - MSR

There are SPSR registers

- hold the saved process state when an exception



16 Confidential © Arm 2018

arm

AArch64 Registers

In AArch64, system configuration is controlled through system registers

System registers are suffixed with “_ELx”, for example SCTLR_EL1

- Suffix defines the lowest exception level that can access that system register
- For example:
 - `TTBRO_EL1` Can be accessed from EL1, EL2, and EL3
 - `TTBRO_EL3` Can be accessed from EL3

Use the MRS instruction to read a system register, and MSR instruction to write to a system register

- `MRS` `X0, SCTLR_EL1` ; `X0 = SCTLR_EL1`
- `MSR` `SCTLR_EL1, X0` ; `SCTLR_EL1 = X0`

Agenda

Arm v8.x overview

Registers

ISA

SVE

Exception

Memory management

Memory model

Atomic

Architecture timer

Virtualization

MPAM

Security

RAS



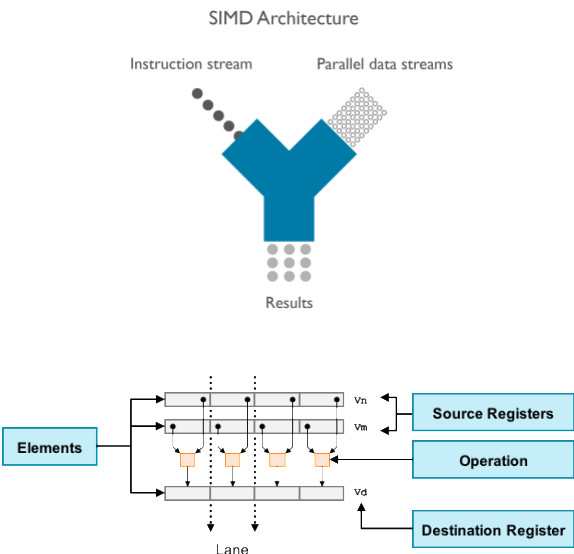
AArch64 Advance SIMD ISA

NEON is a wide SIMD data processing architecture

- Extension of the ARM® instruction set
- Armv8 NEON has 32 registers, 128-bits wide.

NEON Instructions perform “Packed SIMD” processing

- Registers are considered as vectors of elements of the same data type
- Data types can be: signed/unsigned 8-bit, 16-bit, 32-bit, 64-bit, single/double prec. Float (Armv8.2-A also support half prec Float)
- Instructions perform the same operation in all lanes



arm

21 Confidential © Arm 2018

Sample SIMD Instructions - ADD

- **ADD V2.4H, V0.4H, V1.4H**
 - Register split into equal size and type elements (four 16-bit elements)
 - Operation performed on same element of each register

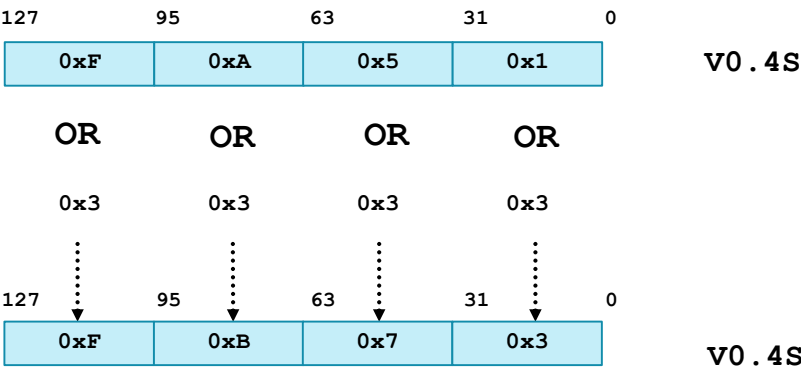
63	47	31	15	0	
0x1001 0x1234 0x7 0xAB					V0.4H
+					
0xFF0 0x5678 0xFFF8 0xCD					V1.4H
=					
0x1FF1 0x68AC 0xFFFF 0x178					V2.4H

arm

22 Confidential © Arm 2018

Sample Instructions - ORR

- ORR V0.4S, #0x3
 - Bitwise inclusive OR each element of V0.4S with constant 0x3



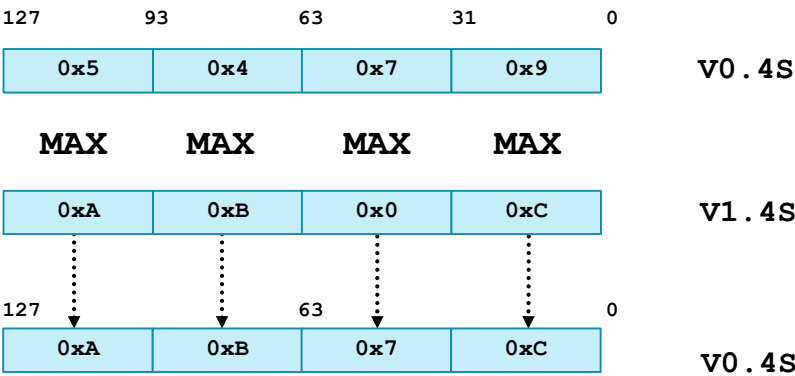
23 Confidential © Arm 2018

arm

Sample Instructions - SMAX

SMAX V0.4S, V0.4S, V1.4S

- Get the signed maximum of the same elements in V0 and V1




24 Confidential © Arm 2018

arm

Sample Instructions - SMULL

- **SMULL** **V0.2D, V1.2S, V2.2S**
 - Signed multiply long - gets 64-bit integer elements from multiplication of 32-bit integer elements




25 Confidential © Arm 2018

arm

Sample Instructions - BIT

- **BIT** **V0.8B, V1.8B, V2.8B**
 - Insert each bit from V1 into V0 if the corresponding bit of V2 is 1, otherwise leave V0 bit unchanged



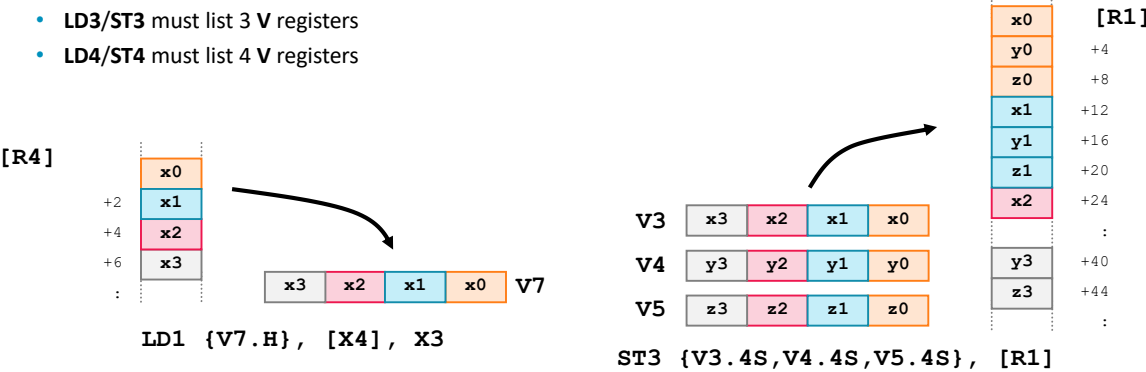
26 Confidential © Arm 2018

arm

Multiple <n> - Element Structure Access

LD<n> and ST<n> can transfer multiple structures between memory and registers

- LD1/ST1 can list 1 – 4 V registers
- LD2/ST2 can list 2 V registers
- LD3/ST3 must list 3 V registers
- LD4/ST4 must list 4 V registers



27 Confidential © Arm 2018

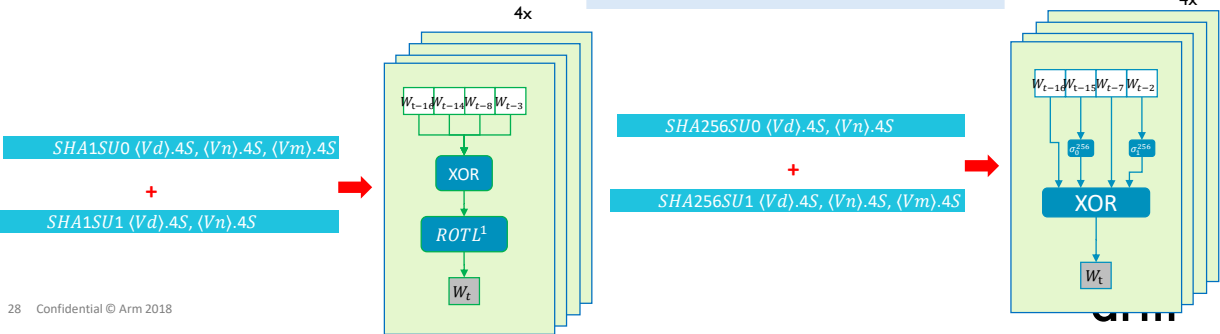
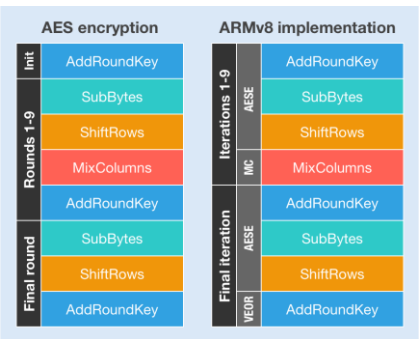
arm

AArch64 Crypto ISA

Crypto extension AES

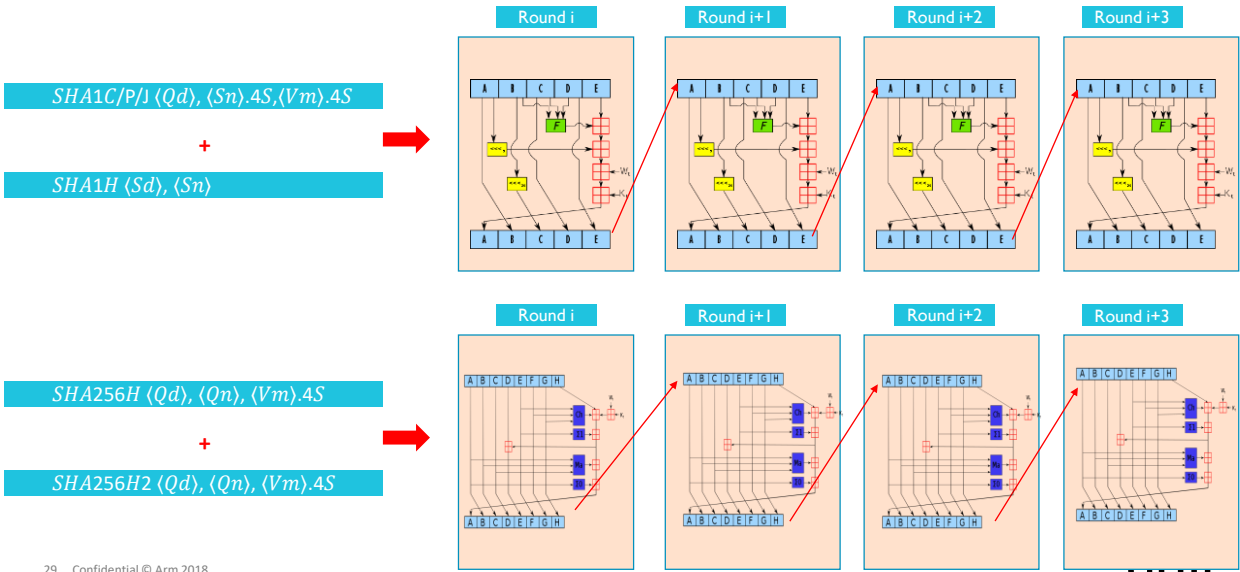
- AESE (AES single round encryption)
- AESD (AES single round decryption)
- AESMC (AES mix columns)
- AESIMC (AES inverse mix columns)

Crypto extension for SHA1 and SHA256



28 Confidential © Arm 2018

AArch64 Crypto ISA



29 Confidential © Arm 2018

arm

ARMv8 CRC32 instructions

CRC instructions operate on the general-purpose register file to update a 32-bit CRC sum from an input value of 8, 16, 32 or 64 bits.

Two different families of CRC instruction are provided to support two commonly used polynomials.

Example instruction	Description
CRC32B Wd, Wn, Wm	Accumulate one byte of input data from Wm<7:0> into the 32-bit CRC sum from Wn, and write the updated sum to Wd. Uses a polynomial of 0x04C11DB7.
CRC32X Wd, Wn, Xm	Accumulate one doubleword (eight bytes) of input data from Xm into the 32-bit CRC sum from Wn, and write the updated sum to Wd. Uses a polynomial of 0x04C11DB7.
CRC32CB Wd, Wn, Wm	Accumulate one byte of input data from Wm<7:0> into the 32-bit CRC sum from Wn, and write the updated sum to Wd. Uses a polynomial of 0x1EDC6F41.
CRC32CX Wd, Wn, Xm	Accumulate one doubleword (eight bytes) of input data from Xm into the 32-bit CRC sum from Wn, and write the updated sum to Wd. Uses a polynomial of 0x1EDC6F41.

```
static u32 crc32_arm64_le_hw(u32 crc, const u8 *p, unsigned int len)
{
    s64 length = len;

    while ((length -= sizeof(u64)) >= 0) {
        CRC32X(crc, get_unaligned_le64(p));
        p += sizeof(u64);
    }

    /* The following is more efficient than the straight loop */
    if (length & sizeof(u32)) {
        CRC32u(crc, get_unaligned_le32(p));
        p += sizeof(u32);
    }
    if (length & sizeof(u16)) {
        CRC32h(crc, get_unaligned_le16(p));
        p += sizeof(u16);
    }
    if (length & sizeof(u8))
        CRC32B(crc, *p);

    return crc;
}
```

30 Confidential © Arm 2018

arm

Agenda

- Armv8.x overview
- Registers
- ISA
- SVE
- Exception
- Memory management
- Memory model
- Atomic
- Architecture timer
- Virtualization
- MPAM
- Security
- RAS

31 Confidential © Arm 2018



SVE

A vector extension to the Armv8-A architecture with some major new features:



Gather-load and scatter-store

Loads a single register from several non-contiguous memory locations.

	1	2	3	4
+	5	5	5	5
pred	1	0	1	0
=	6	2	8	4


Per-lane predication

Operations work on individual lanes under control of a predicate register.

```
for (i = 0; i < n; ++i)
INDEX i  n-2  n-1  n  n+1
CMPLT n  1  1  0  0
```

Predicate-driven loop control and management

Eliminate scalar loop heads and tails by processing partial vectors.



Vector partitioning and software-managed speculation

First Faulting Load instructions allow memory accesses to cross into invalid pages.



No preferred vector width

The above features allow the production of compiled binaries that are agnostic to hardware vector width (which can be between 128-2048 bit at 128 bit increments).

32 Confidential © Arm 2018

arm

New SVE architectural state

Scalable vector registers


- Z0-Z31 extending NEON's V0-V31.
- Packed DP, SP & HP floating-point elements.
- Packed 64, 32, 16 & 8-bit integer elements.

Scalable predicate registers

- P0-P7 governing predicates for load/store/arithmetic.
- P8-P15 additional predicates for loop management.
- FFR first fault register for speculation.

Scalable vector control registers

- ZCR_ELx vector length (LEN=1..16).
- For exception/privilege levels EL1 to EL3.



33 Confidential © Arm 2018

arm

Agenda

Armv8.x overview

Registers

ISA

SVE

Exception

Memory management

Memory model

Atomic

Architecture timer

Virtualization

MPAM

Security

RAS



34 Confidential © Arm 2018

AArch64 exceptions

In AArch64, exceptions are split between:


- Synchronous: Data Aborts from the MMU, Permission Faults, Alignment Faults, service call instructions (e.g. **SVC**), etc
- Asynchronous: IRQs, FIQs, SErrors (System Errors)

On taking an exception, the EL can either stay the same or get higher

- Exceptions are never taken to EL0

Asynchronous exceptions can be routed to a higher EL


- **HCR_EL2** controls routing to EL2
- **SCR_EL3** controls routing to EL3
- Separate bits to control routing of IRQs, FIQs, and SErrors



35 Confidential © Arm 2018

arm

Taking an exception



When an exception occurs:

- **SPSR_ELx** updated
- **PSTATE** updated (EL stays the same or gets higher)
- Return address stored to **ELR_ELx**
- PC set to vector address
- **ESR_ELx** updated with cause of exception
 - Only if synchronous or SError exception

Execute an **ERET** instruction to return from an exception:

- Restores **PSTATE** from **SPSR_ELx**
- Restores PC from **ELR_ELx**

36 Confidential © Arm 2018

arm

Agenda

Armv8.x overview

Registers

ISA

SVE

Exception

Memory management

Memory model

Atomic

Architecture timer

Virtualization

MPAM

Security


RAS

37 Confidential © Arm 2018



What is a Memory Management Unit?

- **The MMU handles the translation of virtual address to physical addresses**
 - Provides hardware to read translation tables in memory (“table walking”)
 - Translation Table Base Registers (**TTBR**) hold the physical base address of the tables
 - Translation Lookaside Buffers (TLBs) cache recent translations
- **When the MMU is enabled, all accesses made by the core are passed through it**
 - MMU will use cached translations from the TLB(s), or perform a table walk
 - Translation must occur before a cache lookup can complete



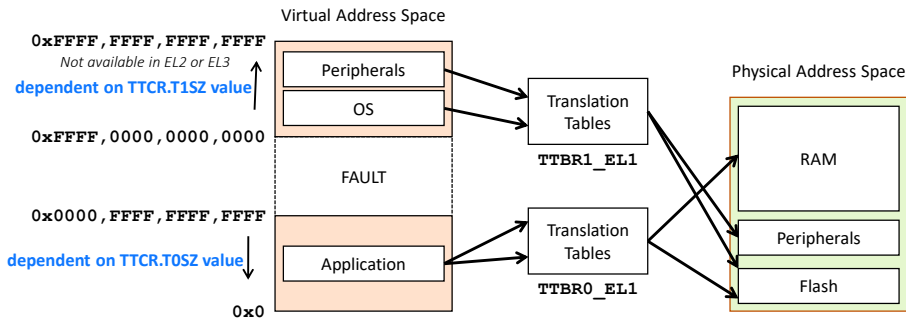
38 Confidential © Arm 2018

arm

Virtual address space

Virtual addresses are 64-bit wide, but not all addresses are accessible

- Virtual memory address space split between two translation tables
 - Each covering a configurable size, up to 48 bits of address space (TCR_ELx)
- Addresses not covered by either translation table automatically generate translation faults



39 Confidential © Arm 2018

arm

Multiple virtual address spaces

A system may define multiple virtual address spaces

OS / Applications

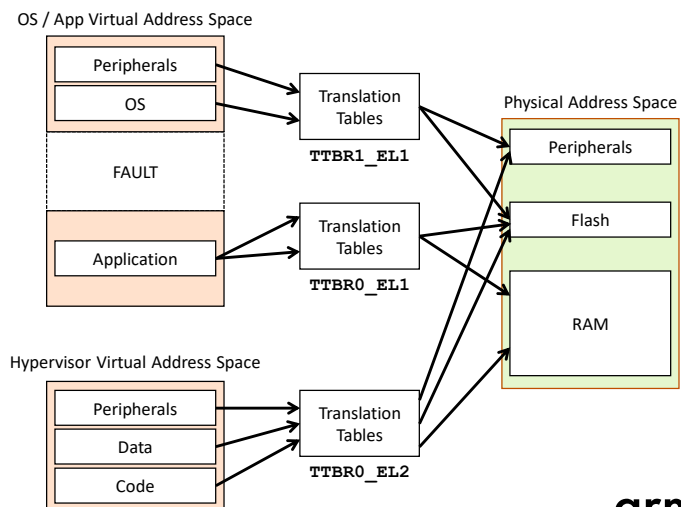
- TTBR0_EL1
- TTBR1_EL1
- TCR_EL1

Hypervisor

- TTBR0_EL2
- TCR_EL2

Secure Monitor

- TTBR0_EL3
- TCR_EL3



40 Confidential © Arm 2018

arm

AArch64 Translation Tables

AArch64 supports 3 different translation granules

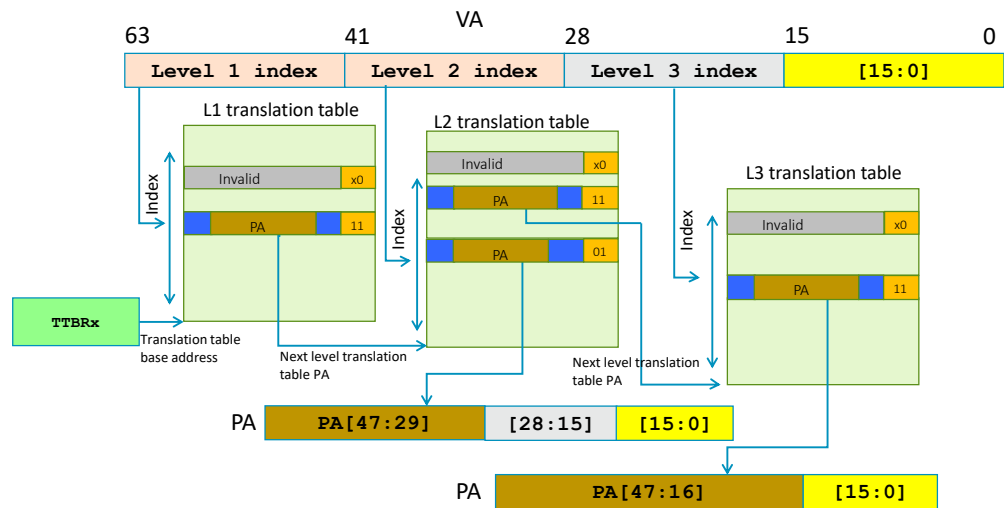
- 4KB, 16KB, or 64KB
 - Defines block size at lowest level of translation table and size of tables
- Configurable for each TTBR
- It is IMPLEMENTATION DEFINED which of the three are supported
 - ID_AA64MMFR0_EL1 reports supported sizes

Larger granules reduce the number of levels of table required

- Particularly important in Virtualized systems

Translation table and address translation

64KB Granule translation example



Agenda

- Armv8.x overview
- Registers
- ISA
- SVE
- Exception
- Memory management
- Memory model
- Atomic
- Architecture timer
- Virtualization
- MPAM
- Security
- RAS

43 Confidential © Arm 2018



Weak memory mode and memory types

The ARM architecture defines a weak ordering model

- This means that accesses *might not* occur in program order
- ARM memory consistency model allows:

Memory access to different locations	Reordered?
Loads reorders after loads	Yes
Loads reorders after stores	Yes
Stores reorders after stores	Yes
Loads reorders after stores	Yes
Atomic reordered with loads	Yes
Atomic reordered with stores	Yes
Dependent loads reordered	No

ARM architecture defines memory types to tell the processor how it can access that location

- Normal memory
- Device memory

44 Confidential © Arm 2018

arm

Memory types

Normal

- Used for code and data
- Processor allowed to re-order, re-size and repeat accesses
- Speculative accesses allowed

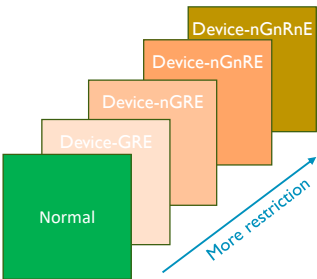
Device

- Used for peripherals
- Accesses could have side effects, so there are more restrictions on what optimizations a processor can perform
- Speculative data accesses not allowed

Other attributes can also be specified

- For example whether a region is executable, shareable, and cacheable

Speculative instruction fetches are allowed to any region that's executable at *some* exception level



Memory Types: Normal

The Normal type is used for code and most data regions

Normal memory gives the best performance because it imposes the fewest restrictions


- Allows the processor to re-order, repeat, and merge accesses

For optimal performance, application code and data should be marked as Normal

- Ordering can still be enforced when required using explicit barrier operations

Address regions marked as Normal can be accessed **speculatively**

- Data or instructions fetched from memory before being explicitly referenced



Memory Types: Device


The Device type is used for regions where accesses can have side-effects

- Example: Writing to a peripheral's control register may trigger an interrupt as a side-effect
- Typically only used for peripherals

Device type imposes more restrictions on the core

Attempting to execute from a region marked as Device is UNPREDICTABLE

Speculative data accesses cannot be performed to Device regions




47 Confidential © Arm 2018

arm

Memory Types: Device (2)

Four variants of Device are available:

- Device-nGnRnE - most restrictive
- Device-nGnRE
- Device-nGRE
- Device-GRE - least restrictive



Gathering (G, nG)

- Determines whether multiple accesses can be merged into a single bus transaction
- nG: number/size of accesses on the bus = number/size of accesses in code

Re-ordering (R, nR)

- Determines whether accesses to same device can be re-ordered
- nR: accesses to the same IMPLEMENTATION DEFINED block size will appear on the bus in program order

Early Write Acknowledgement (E, nE)

- Indicates to the memory system whether a buffer can send acknowledgements

48 Confidential © Arm 2018

nE: The response should come from the end slave, not buffering in interconnect

arm

Barriers

The ARM architecture includes barrier instructions to force access order and access completion at a specific point

- **DMB** – Data Memory Barrier
- **DSB** – Data Synchronization Barrier
- **ISB** – Instruction Synchronization Barrier

DMB

Explicit memory accesses before the **DMB** are observed before any explicit access after the **DMB**

- Does not guarantee when the operations happen, just the order

LDR X0, [X1]	← Must be seen by memory system before STR
DMB SY	
ADD X2, #1	← May be executed before or after memory system sees LDR
STR X3, [X4]	← Must be seen by memory system after LDR

The effects of any data/unified cache maintenance operations *issued by this core* before the **DMB** are observed by explicit data accesses after the **DMB**

- No effect on operations broadcast by other cores

DC CVAC, X5	
LDR X0, [X1]	← Effect of data cache clean <i>might</i> not be seen by this instruction
DMB SY	
LDR X2, [X3]	← Effect of data cache clean <i>will</i> be seen by this instruction

DSB


A **DSB** is more restrictive than a **DMB**

- Use a **DSB** when necessary, but do not overuse them

No instruction after a **DSB** will execute until:

- All explicit memory accesses before the **DSB** in program order have completed
- Any outstanding cache/TLB/branch predictor operations complete

DC	ISW	← Operation must have completed before DSB can complete
STR	X0, [X1]	← Access must have completed before DSB can complete
DSB	SY	
ADD	X2, X2, #3	← Cannot be executed until DSB completes



In a multi-core system, if cache/TLB/branch maintenance prediction operation is broadcast – the operation must have completed on all cores that received it


- Operations received by the core via broadcast do not affect **DSBs**

“One-Way” Barriers (1)

AArch64 adds new load/store instructions with implicit barrier semantics

Load-Acquire (**LDAR**)

- All accesses after the **LDAR** are observed after the **LDAR**
- Accesses before the **LDAR** are not affected




Accesses can cross a barrier in one direction but not the other

“One-Way” Barriers (2)

AArch64 adds new load/store instructions with implicit barrier semantics

Store-Release (STLR)

- All accesses before the **STLR** are observed before the **STLR**
- Accesses after the **STLR** are not affected



53 Confidential © Arm 2018

arm


“One-Way” Barriers (3)

LDAR and **STLR** may be used as a pair

- To protect a critical section of code
- May have lower performance impact than a full **DMB**
- No ordering is enforced *within* the critical section

Exclusive versions also available

- **LDAXR, STLXR**
- Remove the need for explicit barrier instructions in synchronization code



54 Confidential © Arm 2018

arm

Agenda

- Armv8.x overview
- Registers
- ISA
- SVE
- Exception
- Memory management
- Memory model
- Atomic
- Architecture timer
- Virtualization
- MPAM
- Security
- RAS





55 Confidential © Arm 2018

Atomic

load/store exclusive

- Split the operation of atomically updating memory into two separate steps, Load exclusive and store exclusive.
- Together, they provide atomic updates in conjunction with exclusive monitors that track exclusive memory accesses


```
l: ldrex w0, [var]
  add w0, w0, #num
  strex w1, w0, [var]
  teq w1, #0
  bne lb
```



Atomic operations in Armv8.1-a onwards with Large System Extension

- add, bitwise ops, max/min in variants with and without returning a result
- swap, compare and swap

```
LDADD w1, w2, [x0]
```



56 Confidential © Arm 2018

Agenda

Armv8.x overview
Registers
ISA
SVE
Exception
Memory management
Memory model
Atomic
Architecture timer
Virtualization
MPAM
Security
RAS

57 Confidential © Arm 2018



Architecture timer

System counter generates a fixed frequency incrementing count, which is distributed to all cores

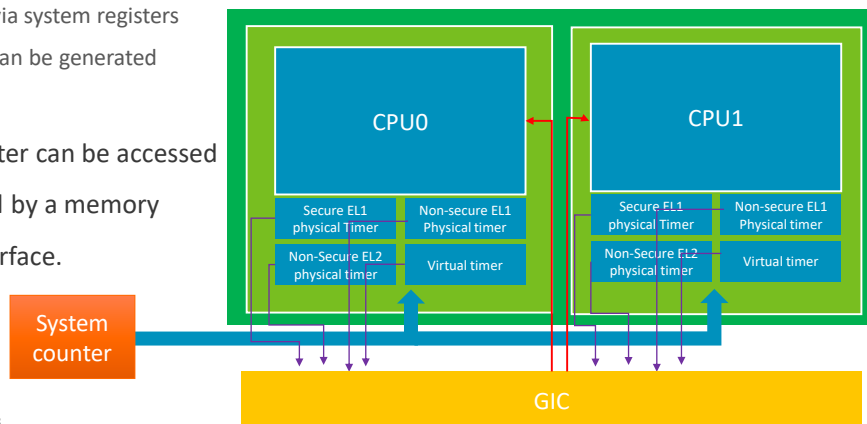
Each core implements a number of comparators

- Secure physical EL1, Non-secure physical EL2, Non-secure EL1 physical and Virtual comparators
- Accessed via system registers
- Interrupt can be generated

System counter can be accessed

Or controlled by a memory

Mapped interface.



58 Confidential © Arm 2018

arm

Agenda

Armv8.x overview

Registers

ISA

SVE

Exception

Memory management

Memory model

Atomic

Architecture timer

Virtualization

MPAM

Security

RAS



59 Confidential © Arm 2018



Virtualization Extension

Arm virtualization Extension provides HW-assisted type1/2 hypervisor support

- CPU virtualization
 - Addition EL for hypervisor.
 - Sensitive instructions trap
 - Hypervisor call instruction
- Memory virtualization
 - Stage 2 translation, IPA->PA
- IO virtualization
 - Stage 2 translation for MMIO
 - Virtual interrupt
 - Physical interrupt can be trapped.
 - Hypervisor software can generate virtual interrupt.
 - GIC virtualization support.
 - Timer virtualization



60 Confidential © Arm 2018

arm

Virtualization Extension- CPU Virtualization

The Hypervisor can be configured to trap certain instructions

- Configured through Hypervisor Control Register (HCR_EL2)
- When trapped, Exception Syndrome Register (ESR_EL2) provides information about trapped instruction

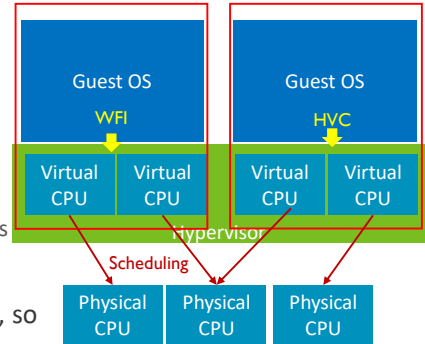
Instructions could be trapped

- System instructions, e.g. cache and TLB maintenance instructions
- Accesses to Auxiliary Control Register (ACTLR_EL1) Reads to ID registers
- WFE and WFI instructions

Some registers have dedicated control for virtualization purpose, so no trapping is required

- E.g., MIDR_EL1, MPIDR_EL1


Hypervisor calls -HVC, guest OS could call services provided by hypervisor



61 Confidential © Arm 2018

arm

Virtual Memory in two Stages

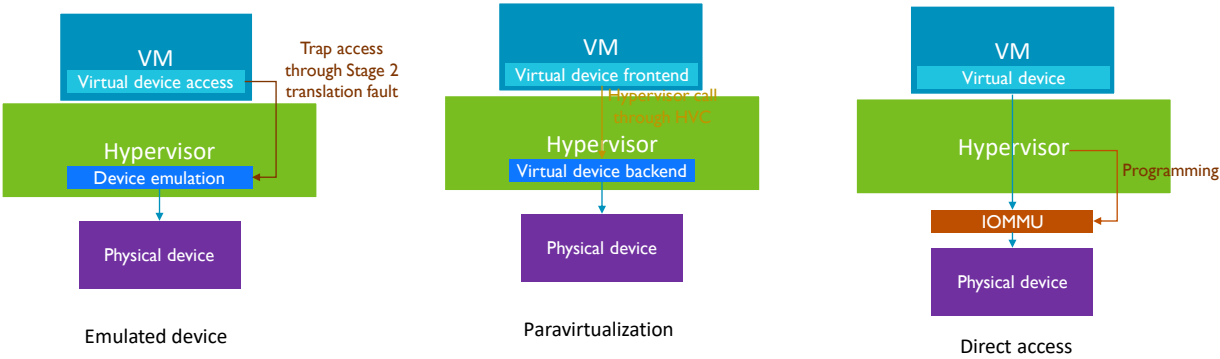


62 Confidential © Arm 2018

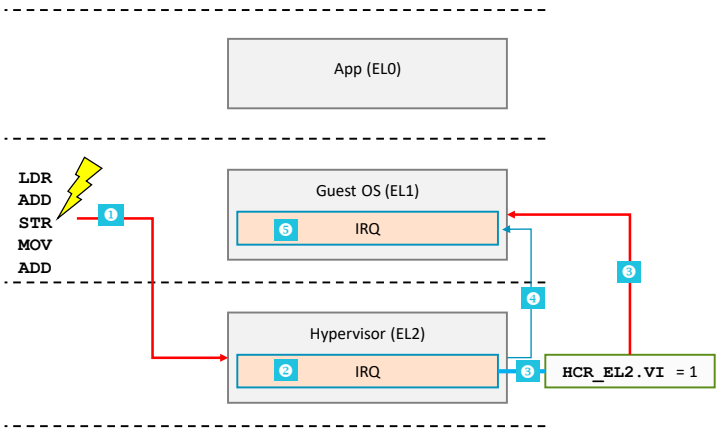
arm

Virtualization Extension- Device virtualization

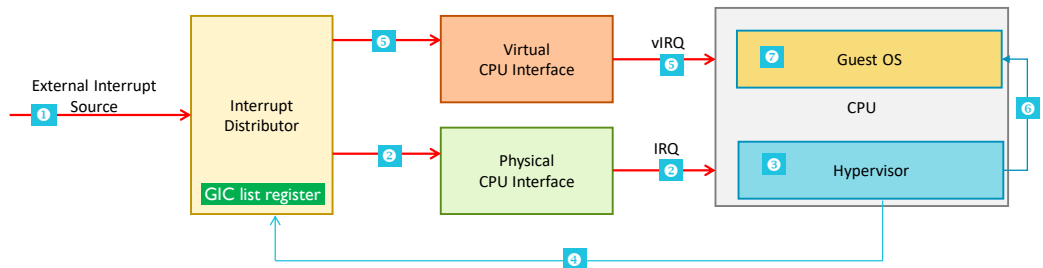
3 methods for device access virtualization



Virtual interrupt signaling (Internal)




Virtual interrupt signaling (GIC)




65 Confidential © Arm 2018

arm

Xen and KVM implementation



Xen – Type1 Hypervisor



KVM – Type2 Hypervisor

66 Confidential © Arm 2018

arm

Overhead without Virtualization Host Extension(VHE)

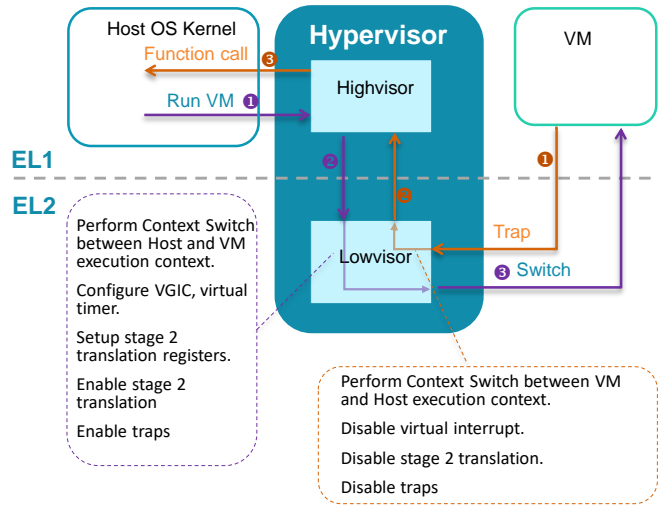
Host OS must run in both EL1 and EL2

- Requires kernel code to be aware where it is running

Calls to EL2 functions must be via HVC trap

- Required for any access to EL2 registers
- Configuring Stage 2 translations

Moving from Guest to Host OS requires a full context switch of EL1 system registers



Pre-VHE KVM Virtualization

67 Confidential © Arm 2018

arm


Armv8.2 Virtualization Host Extension(VHE)

VHE reduces the overhead of transaction between VM and hypervisor

- Allows host kernel and hypervisor totally run in the same EL
- Minimum the need of host kernel modification to support VHE

When software enables VHE (HCR_EL2.E2H=1)

- EL2 translation regime supports two translation table TTBR0/1
- When $HCR_EL2.TGE==1$, translation for EL0 (running host application) uses EL0/EL2 translation regime, stage 2 translation is disabled. When $HCR_EL2.TGE==0$, translation for EL0 (running guest application) uses EL0/EL1 translation regime.
- Add ASID and VMID support in EL2.
- Add virtual timer support in EL2.
- Redirect TTBR0/1_EL1, SCTRL_EL1 etc registers access to EL2 physical registers.
- Introduce new XXX_EL12 (e.g TTBR0_EL12) registers in EL2.




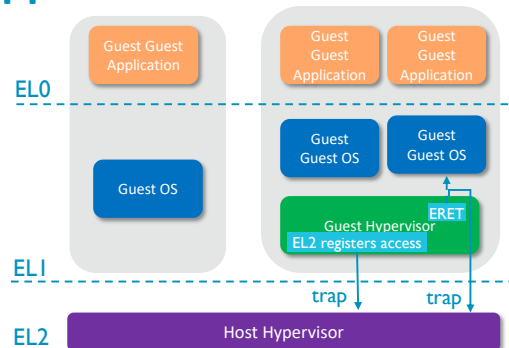
68 Confidential © Arm 2018

arm

Armv8.3 Nested virtualization support

Supports running the guest hypervisor in EL1, controlled by HCR_EL2.NV

- Traps EL2 operations executed in EL1 to EL2
 - Emulating EL2 in EL1
 - shadow EL1 register state that reflects the EL2 register state
- Traps ERET to EL2



Memory Virtualization

- Software combines stage 2 translation of guest hypervisor and stage 2 translation of host hypervisor to a shadow translation table

arm

Armv8.4 Nested virtualization support

Nested Virtualization Extensions (NEVE)



- Provides two techniques to avoid traps based on register classification

VM Registers: Redirection to Memory

- NEVE redirects VM register access instructions to memory
- On nested VM entry, the host hypervisor can get VM register states from memory

Hypervisor Control Registers

- Traps are handled by redirecting to EL1 registers in software



arm

Agenda

- Armv8.x overview
- Registers
- ISA
- SVE
- Exception
- Memory management
- Memory model
- Atomic
- Architecture timer
- Virtualization
- MPAM
- Security
- RAS

71 Confidential © Arm 2018



MPAM (armv8.4-a)

MPAM = Memory system performance resource Partitioning And Monitoring

Software environments labelled with Partition ID (PARTID) and Perf Monitoring Group (PMG)

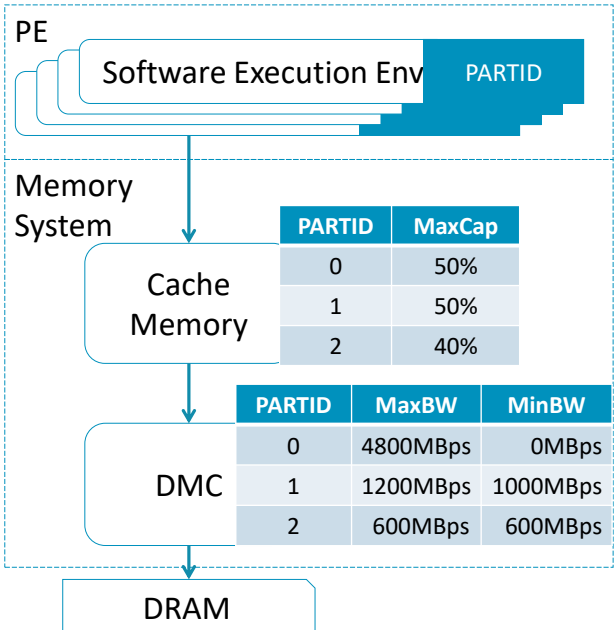
- e.g enabling hypervisors to monitor and control VM memory bandwidth.

PE system registers and behaviors to generate PARTID and PMG to label memory system requests

Propagation through and storage of MPAM PARTID and PMG in interconnects, buffers and caches

- Resource controls within memory system components responsive to PARTID
- Monitors to measure the resources controlled by MPAM

72 Confidential © Arm 2018



arm

Agenda

Armv8.x overview

Registers

ISA

SVE

Exception

Memory management

Memory model

Atomic

Architecture timer

Virtualization

MPAM

Security

RAS

73 Confidential © Arm 2018



TrustZone solution

TrustZone processors have two security states

- Separate MMU configurations
- Registers protection


Memory system is security aware

- Addresses have a security state on the bus
- Enforced by hardware

Peripheral access can be arbitrated by security state

- Trusted code can have access to IO peripherals

Memory system can be implemented to allow dynamic partitioning

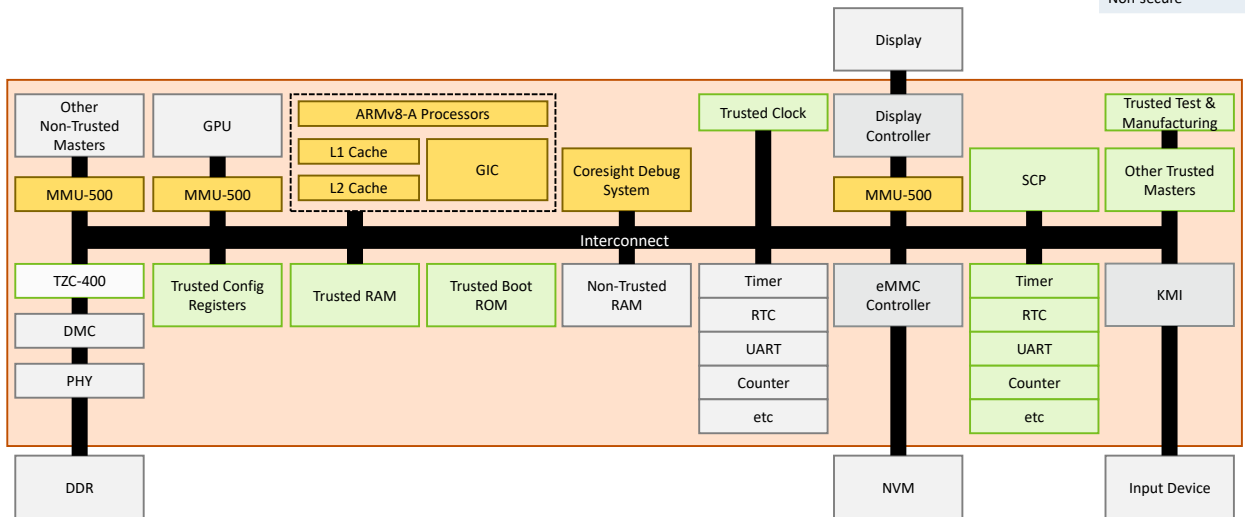


74 Confidential © Arm 2018

74

arm

Example TBSA System



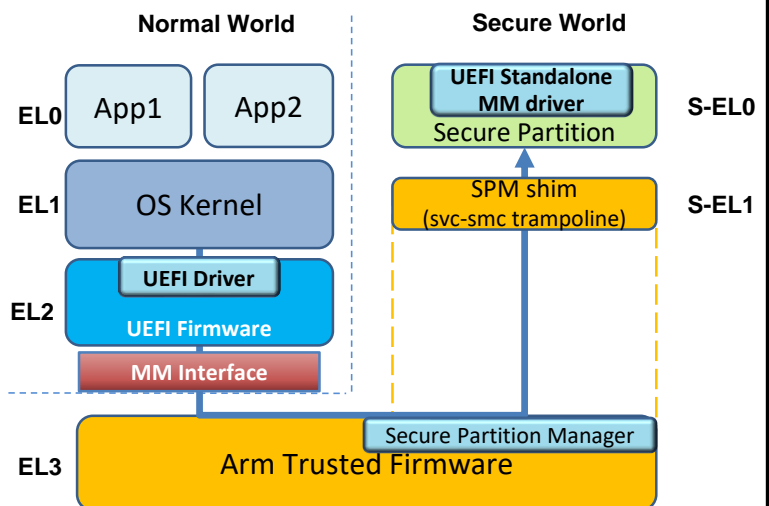
77 Confidential © Arm 2018

arm

Server security

Main design goals:

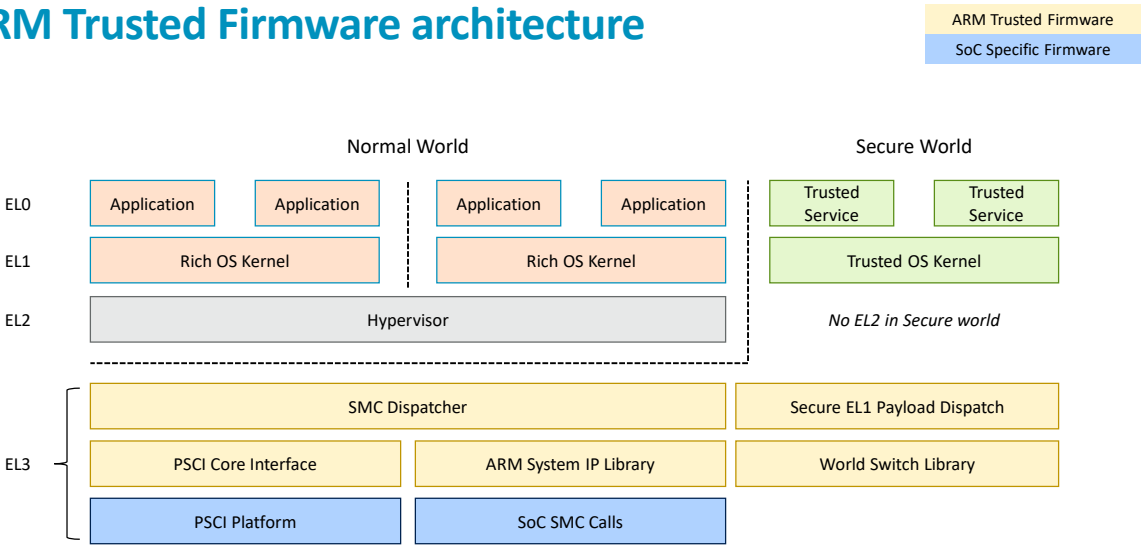
- Isolated execution context
- Limited access to system resources
- OS agnostic
- Leverage Arm MM Interface Specification, Arm TrustZone & UEFI Standalone MM
- Well defined interfaces
- Code reuse between normal/secure world whenever possible
- Reduced services code into privileged firmware (EL3)



78 Confidential © Arm 2018

arm

ARM Trusted Firmware architecture



79 Confidential © Arm 2018

arm

Agenda

- Armv8.x overview
- Registers
- ISA
- SVE
- Exception
- Memory management
- Memory model
- Atomic
- Architecture timer
- Virtualization
- MPAM
- Security
- RAS

80 Confidential © Arm 2018



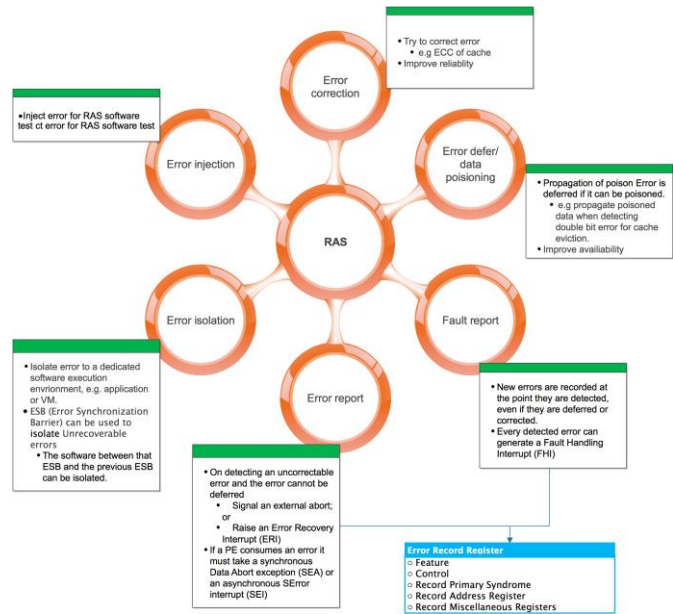
Armv8.2 RAS Extensions

RAS is required for ARMv8.2, it stands for:

- **Reliability** - continuity of correct service
- **Availability** - readiness for correct service
- **Serviceability** - ability to undergo modifications and repairs

RAS extensions provides framework for building RAS features around

- It describes error types, fault/error reporting and handling, and unrecoverable error isolation etc.
- Does not specify RAS features in implementations



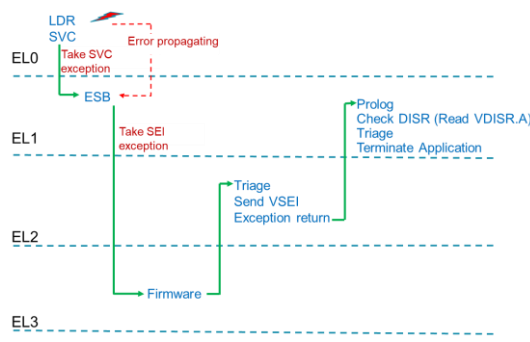
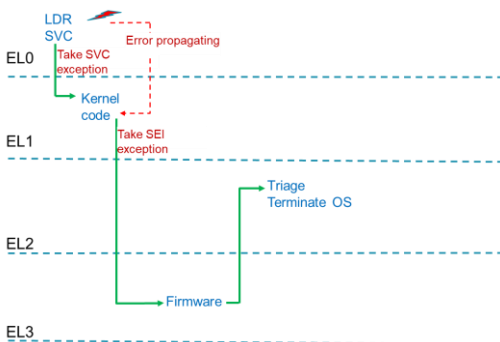
81 Confidential © Arm 2018

arm

Armv8.2 RAS Extensions

ESB is for error isolation between ELs

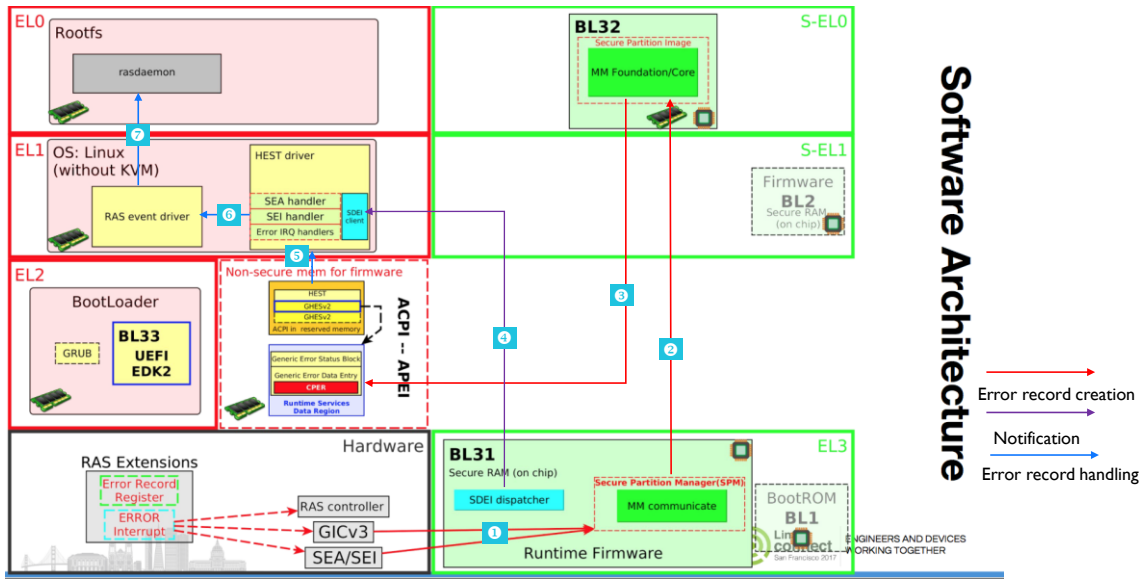
- It guarantees that all Unrecoverable errors generated before the ESB have pended an SEI exception.



82 Confidential © Arm 2018

arm

RAS software



83 Confidential © Arm 2018

Source: Linaro Connect SFO17


arm

arm

The Arm trademarks featured in this presentation are registered trademarks or trademarks of Arm Limited (or its subsidiaries) in the US and/or elsewhere. All rights reserved. All other marks featured may be trademarks of their respective owners.

www.arm.com/company/policies/trademarks

84 Confidential © 2017 Arm Limited



Arm server training

GIC, SMMU

Zenon Xiu
Senior Applications Engineer Manager
Partner Enablement Group

Taipei, July 2018

Copyright © 2018 Arm Limited

What PEG does --Arm Services

Support Services

Training

- Expand employee's knowledge and capabilities
- Reduce time to market
- Fully exploit product features

Technical Support

- Worldwide support
- Design efficiently and effectively
- Errata updates and delivery

Documentation

- Accessible product information
- Comprehensive and up-to-date
- Supplementary online resources

Arm Design Reviews

- End-to-end project assistance
- Optimize subsystem designs
- Reduce risk

www.arm.com/support



1

ARM Training Overview

ARM Training gets partner teams up to speed quickly

- Detailed training across all ARM Products
- Covers advanced technologies such as ARM® DynamIQ and ARM® TrustZone®

Delivered by senior Arm engineers in a way that works best for you and your team



Face to face private training delivered by experienced ARM trainers at a location of your choice. Private courses can be customizable to meet your specific needs.



Live virtual classroom delivered by an experienced ARM trainer. Private and public courses available. Delivered to you, wherever you are at a scheduled date and time.



Online learning - short topics accessible wherever you are and whenever you need them. To be launched Oct 2017.

3 Confidential © Arm 2018

arm

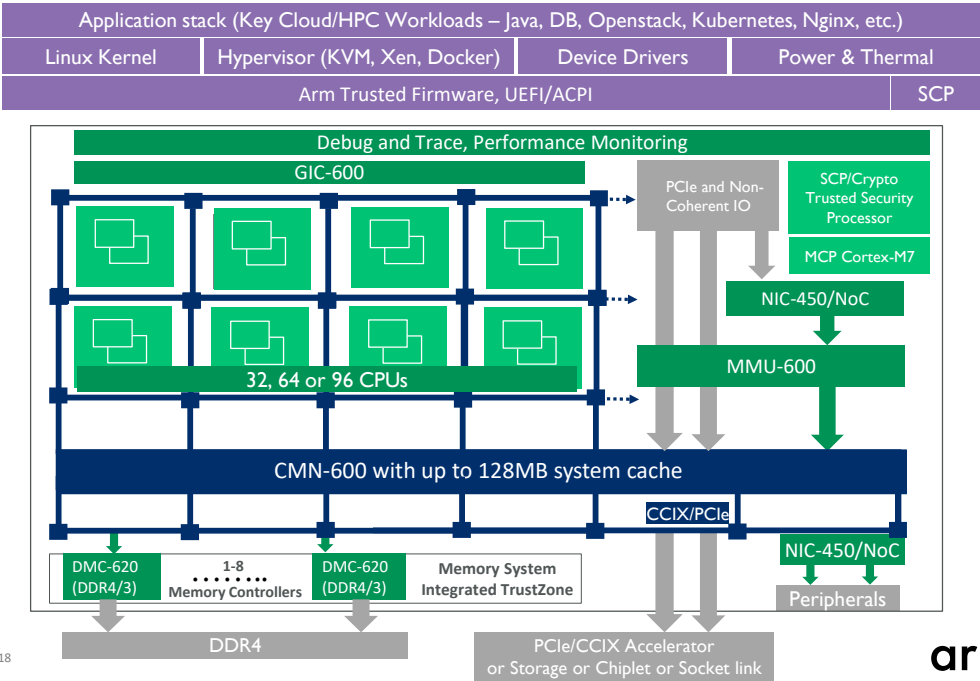
Arm Services

Providing the help you need, when you need it.

A range of **support services** are available to help reduce risk and reduce your time to market.



www.arm.com/support



arm

Agenda

GICv3/v4
SMMUv3



GIC versions

GICv2

Adds:

- Support for virtualization
- Improved handling of Group 1 interrupts by secure software

Implemented by:

- Cortex-A15 MPCore*
- Cortex-A7 MPCore*
- CoreLink GIC-400

* Inclusion of GIC is optional

GICv3

Adds:

- Support for more than 8 cores
- Support for message signalled interrupts
- System Register access to some registers
- Vastly expanded the Interrupt ID space

Implemented by:

- CoreLink GIC-500
- CoreLink GIC-600

GICv4

Adds:

- Direct injection of virtual interrupts

GICv3/v4 overview

0 .. 15: Software Generated Interrupts (SGIs)

- Used for software inter-processor interrupts
- Generated by software to a set of target processors

16 .. 31: Peripheral Interrupts (PPIs)

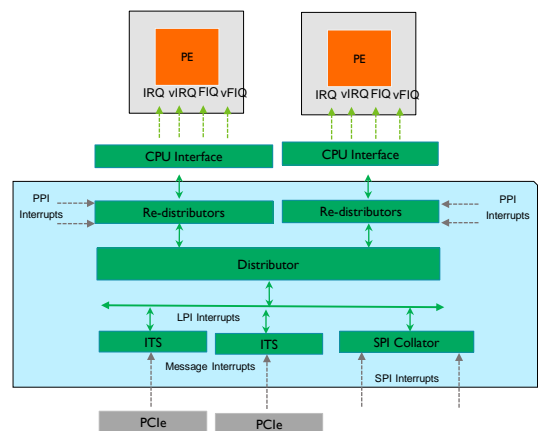
- Banked per processor
- Used for interrupts from peripheral private to a processor

32 .. 991(configurable): Shared Peripheral Interrupts (SPIs)

- Global interrupts that that can be sent to any one processor
 - Either “1 of all” processors or a targeted processor

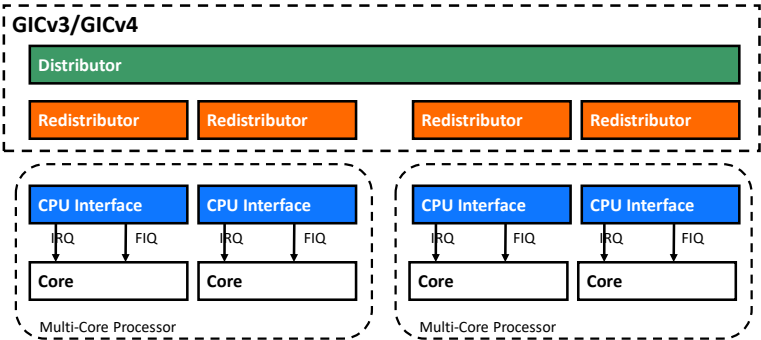
8192+: Locality-Specific Interrupts (LPIs) (*New in GICv3*)

- Interrupts that can be sent to any one targeted processor



Redistributors

GICv3 introduces Redistributors



There is a Redistributor per-connected core, holding settings for private interrupts (PPIs and SGIs)

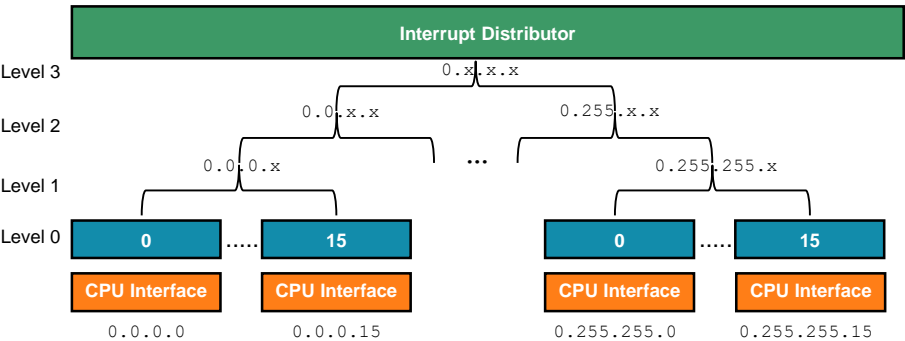
- Allows for distributed designs with Redistributors kept close to target core

Affinity levels and routing

There is increasingly demand for systems with higher core counts

GICv3 increases the number of cores that can be connected

- Connected cores identified by an affinity value, matching system used in ARMv8-A



Interrupt Routing

GICv3 routes interrupt with a “Routing” register

- Routing mode + a 32 bit affinity value
- Affinity corresponds to processor MPIDR

SPI routing (programmed into per interrupt “Routing” register)

- Two routing modes only:
 - “1 of all”
 - “a.b.c.d”

SGI routing (provided by software on SGI generation)

- Two routing modes only:
 - “all except self”
 - “a.b.c.{target list}” where target list is a set of up to 16 processors

11 Confidential © Arm 2018

arm

Models for handling interrupts

Targeted distribution model

- This model applies to all PPIs and to all LPIs

Targeted list model

- This model applies to SGIs only. Multiple PEs receive the interrupt independently.
 - When a PE acknowledges the interrupt, the interrupt pending state is cleared only for that PE

1 of N model

- This model applies to SPIs only.
- The interrupt is targeted at a specified set of PEs, and is taken on only one PE in that set.
 - The PE that takes the interrupt is selected in an IMPLEMENTATION DEFINED manner.

12 Confidential © Arm 2018

arm

SPI, SGI and PPI configuration

Enable

- Only enabled interrupts can be forwarded to a core
- Disabled interrupts can still become pending

Priority

- Each interrupt has an 8-bit priority associated with it

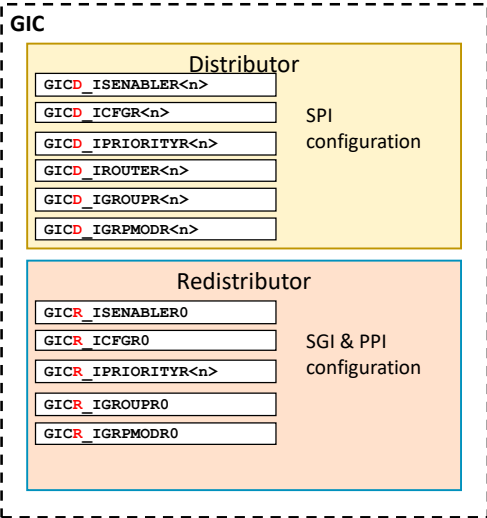
Configuration

- Whether the interrupt is level-sensitive or edge-triggered
- SGIs are always edge-triggered

Routing (SPIs only)

- Interrupt can target one named core, or any core

Security/Group



arm

Introduction to LPIs

LPIs use very large global ID spaces

- State for these ID spaces held in external memory

Physical LPIs

- Each interrupt routed to a single processor at any time
- Configuration table in memory (priorities and enables)
- Per-processor pending table in memory

On GIC-500 & GIC-600, only generated via write to the Interrupt Translation Service

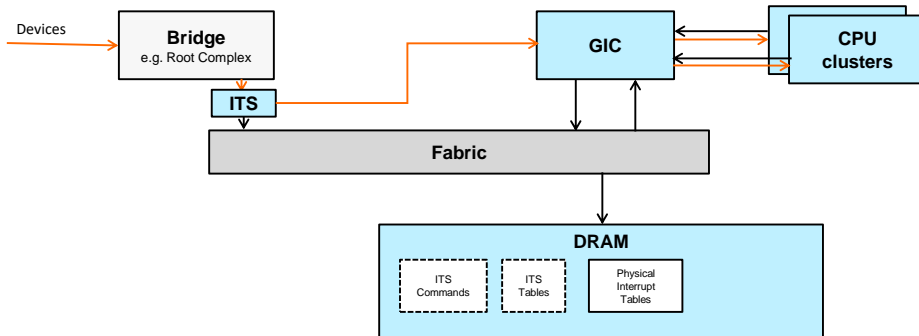
Typically an MSI(-X)

arm

ITS

Message-based interrupts from devices are sent to Interrupt Translation Service in GIC

- Supports standard MSI(-X) interrupts
- Memory-backed to support large number of interrupts
- Cached within GIC to reduce latency and power for common interrupts
- Provides ID translation, allowing devices to be programmed by a VM



15 Confidential © Arm 2018

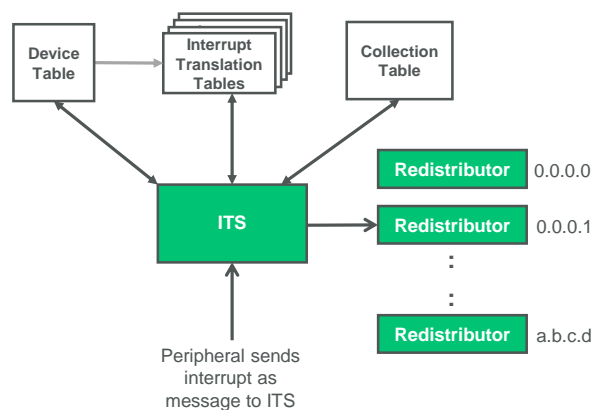
arm

How ITS works

An Interrupt Translation Service (or ITS) maps interrupts to INTIDs and Redistributors

How is an interrupt translated?

- Peripheral sends interrupt as a message to the ITS
 - The message specifies the DeviceID (which peripheral) and an EventID (which interrupt from that peripheral)
- ITS uses the DeviceID to index into the Device Table
 - Returns pointer to a peripheral specific Interrupt Translation Table
- ITS uses the EventID to index into the Interrupt Translation Table
 - Returns the INTID and Collection ID
- ITS uses the Collection ID to index into the Collection Table
 - Returns the target Redistributor
- ITS forwards interrupt to Redistributor



16 Confidential © Arm 2018

arm


ITS - Interrupt Isolation and Translation

Each “device” has a single owner

- Hypervisor and Guest operating system

Each “device” indexes a “Device table” using “Device ID” to select an Interrupt Translation Table (ITT)

- PCIe root complex must drive Requester ID on AW channel e.g. from SMMU



Protection

- Each “device” only be able to generate own interrupts

Translation

- Each “device” can have different mapping: MSI-X ‘data’ => (target CPU, interrupt ID)


17 Confidential © Arm 2018

arm

GICv4

GIC V4 supports the direct injection of virtual interrupts,

- Reduces the hypervisor mediation overhead.
- Map an EventID-DeviceID combination to a vINTID and pINTID(door-bell interrupt) for a specific vPE
 - Optionally a door-bell interrupt can be specified. This is a pINTID that is generated if the vPE is not scheduled when the interrupt is generated.
- Map a vPE to a physical Redistributor



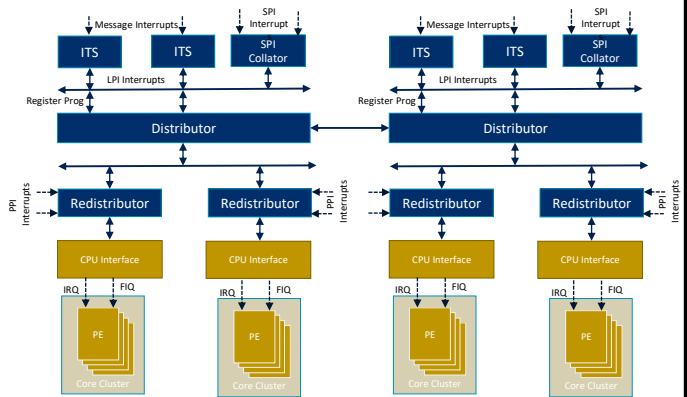
- 1 Hypervisor programs ITS to map EventID/DeviceID to virtual INTID and virtual CPU
- 2 ITS uses Eventide-DeviceID to translation the MSI, and route the interrupt
- 3 If VM is scheduled: Redistributor forwards virtual interrupt vINTID to virtual CPU interface
- 4 If VM is not scheduled: Redistributor forwards door-bell interrupt pINTID to physical CPU interface

18 Confidential © Arm 2018

arm

GIC-600 Features for Infrastructure Solutions

- Affinity Routing (DynamlQ supported format)
- Upto 512 Cores/Threads (128 in GIC-500)
- Coherent Multi-chip
- Multi ITS block support
- Block Interrupt Migration



Key Reasons for GIC-600

1. Multi ITS block support/socket
2. Large Core Count
3. Multi-Chip Support

arm

Agenda

GICv3/v4
SMMUv3



SMMU

System MMU (or SMMU) provides memory management for DMA masters


- CPU MMU only used for the processor cores
- Sometimes referred to as IOMMU

Provides memory translation

- Enables large regions of memory to be allocated without being physically contiguous
- Masters with small addresses (e.g 32-bit) can address all physical memory (e.g 40-bit)

Provides isolation and protection from accessing illegal memory locations

- From process to process or from VM guest to guest
- When using virtualization masters & software do not need to be aware of Hypervisor

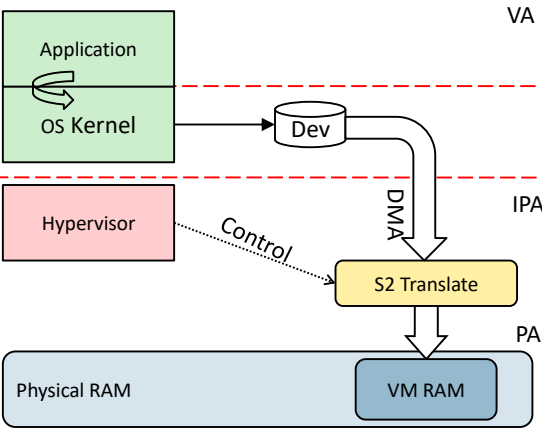


21 Confidential © Arm 2018

arm

SMMU example usage 1

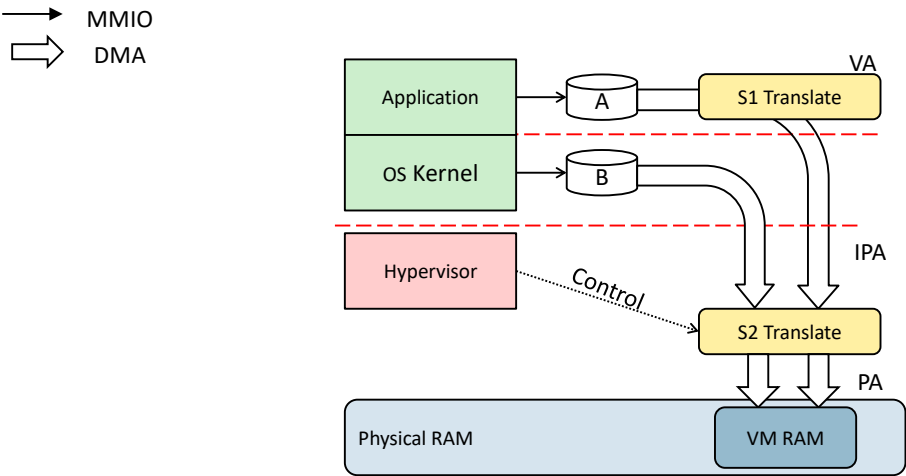
→ MMIO
⇨ DMA



22 Confidential © Arm 2018

arm

SMMU example usage 2



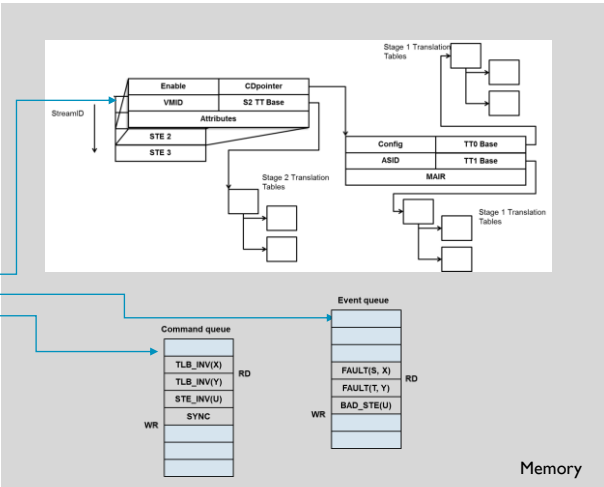
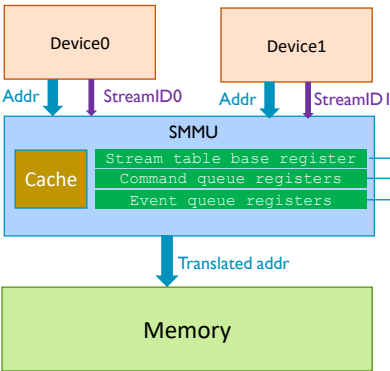
23 Confidential © Arm 2018

arm

SMMUv3

A StreamID is used to determine which context should be used for translation

- A transactions might require both stage 1 and stage 2 translations.

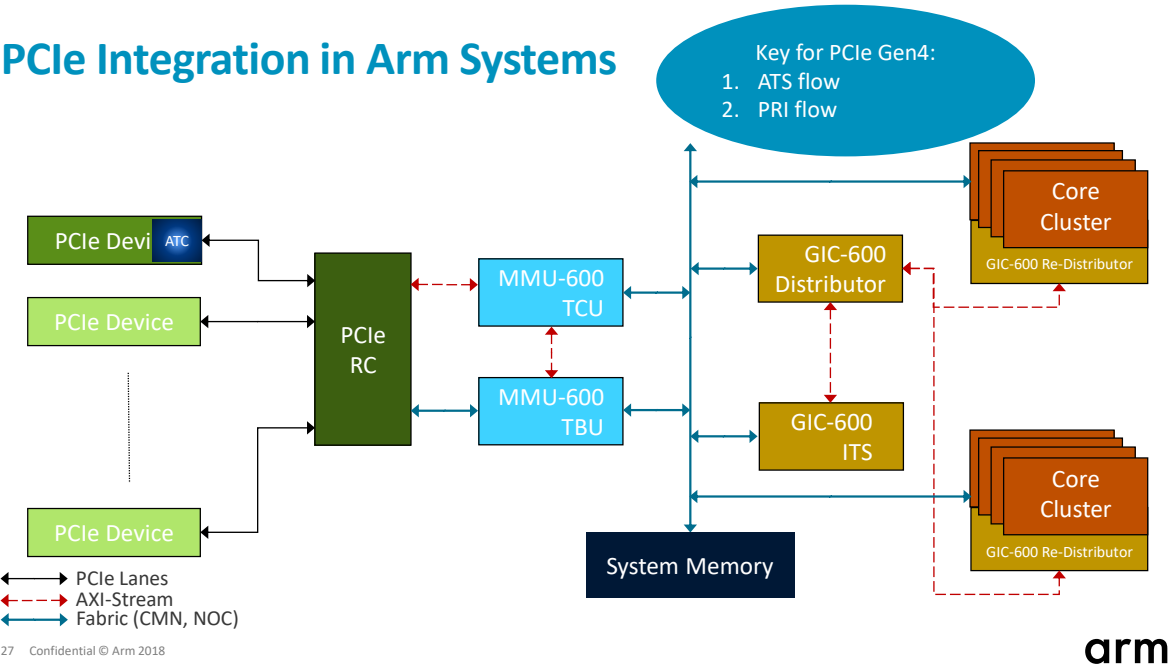


24 Confidential © Arm 2018

SMMU is controlled via a circular command queue in memory
SMMU uses an Event queue to report information back to software

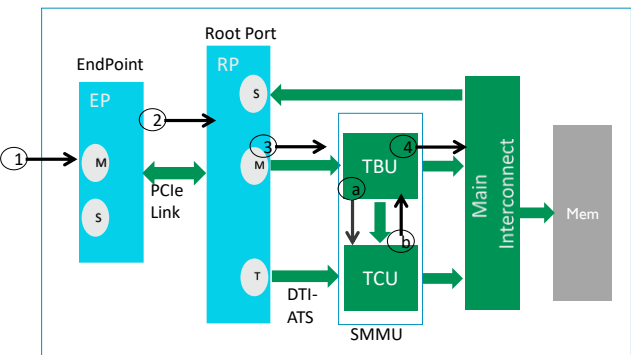
arm

PCIe Integration in Arm Systems



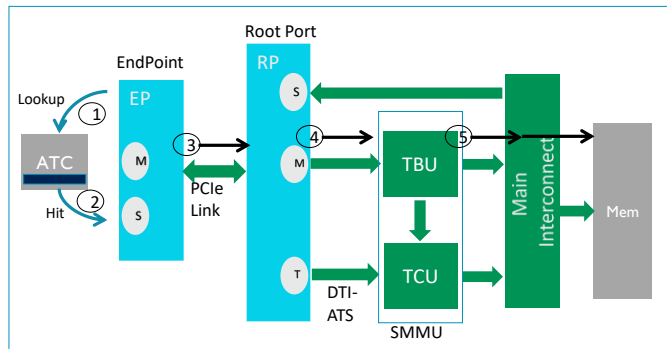
SMMU with PCIe operation – no ATS

1. Client logic generates a TLP with a un-translated address
2. EP sends this as a PCIe TLP to the RP
3. On receipt by the RP, since the packet is a data flow packet, this is sent on the “M” interface
 - a) If the TBU does not have a suitable translation for the address received, it will issue a request to TCU
 - b) The TCU will respond with the response for the TBU
4. The TBU then forwards the transaction to the memory



SMMU with PCIe operation – with ATS and ATC hit

1. Client logic generates a TLP with a virtual address
2. The client logic uses the translated Addr if available from the ATC
3. The EP sends this as a PCIe TLP that has translated address
4. On receipt by the RP, since the packet is a data-flow packet, this is sent on the “M” interface
5. The TBU then forwards the transaction to the memory via the main interconnect



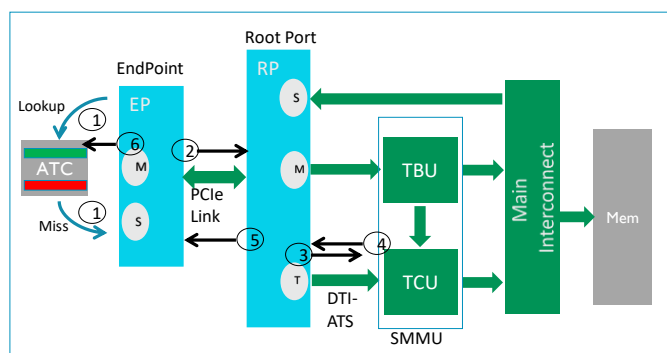
29 Confidential © Arm 2018

29

arm

SMMU with PCIe operation – with ATS and ATC miss

1. EP client generates a PCIe translation for a particular address that needs translation
2. Translation request goes out on the PCIe link to the RP
3. RP sends the translation request it received on the “T” interface to the TCU
4. The TCU then generates the response completion
5. The RP repacks the translation completion TLP back to the EP
6. Once the EP received this completion for the translation request it generated, it populates the local ATC



30 Confidential © Arm 2018

30

arm

The ARM logo, consisting of the lowercase letters 'arm' in a white, sans-serif font, is positioned on the right side of a large blue rectangular area.

The Arm trademarks featured in this presentation are registered trademarks or trademarks of Arm Limited (or its subsidiaries) in the US and/or elsewhere. All rights reserved. All other marks featured may be trademarks of their respective owners.

www.arm.com/company/policies/trademarks